

LongTail Scout — Slide Deck Outline

Submission for **Bright Data Web Data UNLOCKED** hackathon, Track 1 (GTM Intelligence). 2026-05-30. 8 slides, designed for skimming. Paste into Google Slides / Figma / Canva.

This deck is OPTIONAL — lablab.ai doesn't require slides, just the form + video. But if you want supporting material to send to the partner-prize judges (Bright Data, Cognee, Speechmatics), this is the deck.

Slide 1 — LongTail Scout

*** Apollo for the long tail.*

Live AI prospect scout for the 7M small American businesses Apollo can't see.

Bright Data Web Data UNLOCKED hackathon · Track 1, GTM Intelligence · 2026-05-30 <https://longtailscout.com>

Slide 2 — The buyer's problem

The GTM teams at **ServiceTitan, AccuLynx, JobNimbus, Brightwheel, Jobber, Bambee, Toast** spend their weeks chasing the same kind of account:

Small, local, operator-owned, web-first. Real revenue, real budget, no LinkedIn corporate page, no Apollo record.

One SDR builds ~~10 qualified accounts a day by hand — scrolling directory sites, opening every roofing company's homepage, copying URLs out of Google.~~ ****\$300 per account just to FIND them****, before any outreach starts.

Apollo and ZoomInfo don't help here. Their data was built around the LinkedIn employee graph. For the long tail — local trades, niche manufacturers, family-owned services, franchisees — that graph is empty.

This is the job we automate.

Slide 3 — The product

Type a niche × city query. Get a ranked, cited, evidence-backed list of operators in ~40 seconds.

Example query: roofing contractors in Houston

Output: 8-10 ranked operators with

- Company name + homepage URL
- Hiring signals scraped from the careers page (role count + titles)
- Recent activity headlines pulled live from press / blog
- Size estimate from public footprint
- A one-sentence **sales angle** generated from scraped facts
- **Every claim citation-linked to the source URL**

Per-row drill-down shows the underlying Bright Data fetches inline. On-page map plots operators against a heat layer of every other business in the demand index for that niche.

(screenshot here from the live demo)

Slide 4 — How Bright Data powers it

Three-phase hybrid agent pipeline. Bright Data is the spine.

Phase	Step	Bright Data tool
1. Discovery	LLM fires 4 parallel SERP queries: direct, hiring, news, suppliers	Brave SERP API → DDG → BD Scraping Browser cascade. BD handles the long-tail queries Brave misses.
2. Enrichment	Per candidate: render homepage, careers, news. Hybrid plain-fetch + BD.	BD Scraping Browser for JS-heavy ATS pages (Greenhouse / Lever / Workday)

Phase	Step	Bright Data tool
3. Synthesis	DeepSeek ranks + writes citation-grounded sales angles	—

A tiny Playwright bridge service on the home server connects to Bright Data's **WSS (Chrome DevTools Protocol) Browser API** and exposes `/render` + `/serp` + `/screenshot` as HTTP for the Worker.

One BD zone (`lts_browser` , `type browser_api`) handles SERP + render + screenshot. Mutex serialization avoids BD's per-context throttling.

Slide 5 — The killer differentiator: Niche Recon

The feature Apollo physically cannot replicate.

Paste a product description. Get the top 5 long-tail verticals to hunt in.

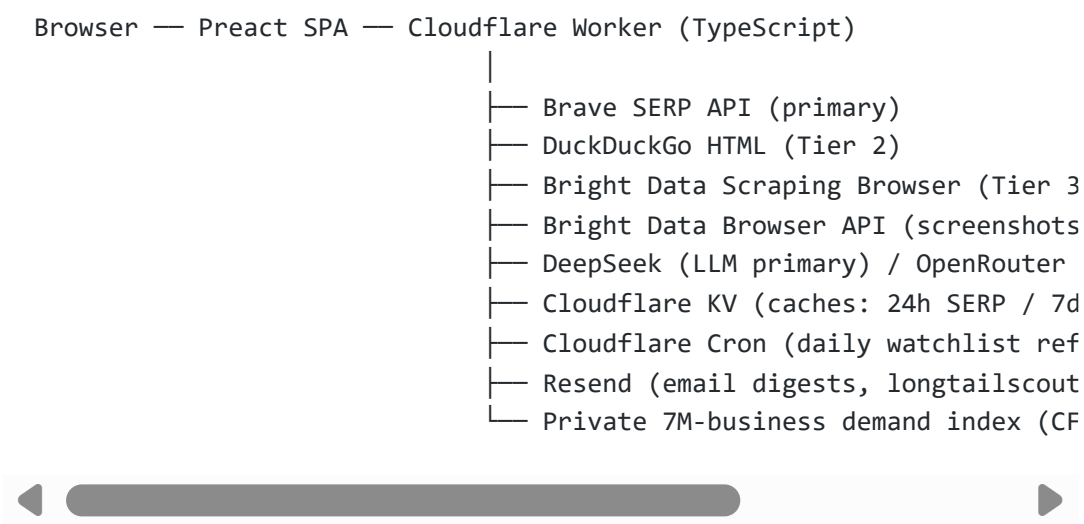
How: we map the description to candidate verticals via LLM, then cross-reference each against our private 7M-business demand index. We rank by **Apollo-thinness** — what share of businesses in the vertical have NO own-domain website. Booking platforms (Booksy, Boulevard, ServiceTitan, Vagaro), social profiles (Facebook, Instagram, Yelp), Google profile pages — all of those make the operator invisible to Apollo's domain-match enrichment.

Example output for "home services CRM for trades" :

Niche	Businesses	Apollo-thin
electrical	30+	100% (every one uses ServiceTitan booking)
hvac	30+	86%
pool service	30+	76%
plumbing	30+	50%

This is the moat. Apollo doesn't have the demand index. Without the index, this feature doesn't exist.

Slide 6 — Architecture



Hosting: Cloudflare Workers (paid plan), custom domain via auto-cert, SSE streaming for the live agent trace. Demo password gates writes; all read paths still require Bearer auth for the demand index (we don't expose the moat).

Stack: Preact + Tailwind v4 + Vite (frontend), Wrangler 3 (deploy), playwright-core + cheerio (bridge), Vitest (118 tests).

Slide 7 — Beyond the demo: MCP, watchlists, security

MCP Server at `/api/mcp` — 12 tools any MCP-aware client (Claude Desktop, Cursor, ChatGPT) can call:

Tool	Use case
scout	Run a full long-tail scout for a niche × city

Tool	Use case
niche_recon	Reverse the funnel — paste a product, get the top verticals
find_businesses	Geotagged businesses from the demand index
demand_count	Single-integer count for a niche
operator_screenshot	Live BD-rendered homepage screenshot
draft_email	Personalized cold email from operator facts
linkedin_check	Apollo-blind verification — is the operator on LinkedIn? (via BD)
find_contacts	Real email / phone / contact from the contact pages (via BD)
decision_maker	Owner/founder + their LinkedIn profile (via BD)
signal_radar	Live third-party news/funding/expansion triggers (via BD)
rank_triggers	Re-rank operators by buying-signal strength
account_brief	One-page Markdown dossier for an operator

Watchlist + alerts: save a query, daily cron refreshes the demand-index count, Resend email digest + Slack/Discord webhook fire when delta > 0. CAN-SPAM compliant — each digest email has an HMAC-signed unsubscribe link.

Security: 118 unit tests covering SSRF guards on both the worker AND the bridge (caught an IPv6 bypass before deploy), HMAC unsubscribe tokens, watchlist PII redaction, the Apollo-thinness platform-host filter, and every new classifier. Two external code reviews (Codex), 10/10 findings closed; full live burn-test of all endpoints + MCP tools.

Slide 8 — What's next + thanks

Next 30 days:

- CRM connectors (HubSpot, Salesforce, Apollo CSV reverse-import)
- Quality mode: opt-in BD-bridge enrichment that recovers the ~35% of candidates plain-fetch drops
- Vertical-tuned prompts (roofing, HVAC, childcare, EMS — different hiring signal weights per market)
- Bright Data MCP Server integration once it's GA — deprecate our Playwright bridge, call BD directly

Thanks to:

- **Bright Data** for the Scraping Browser product and the hackathon promo credit
- **Cloudflare** for Workers, KV, Tunnels, custom-domain auto-cert (the free tier is staggeringly generous)
- **DeepSeek** for cheap fast tool-use LLM — primary all the way
- **lablab.ai** for running the hackathon

Live demo: <https://longtailscout.com> · **Source:** github.com/dholzric/longtail-scout