

# lablab.ai submission — ready-to-paste copy

---

This file has the exact text for every field in the lablab.ai submission form for the **Bright Data Web Data UNLOCKED** hackathon. Copy each block as-is.

**Deadline:** 2026-05-30. **Track:** GTM Intelligence (Track 1). **Required:** demonstrably use  $\geq 1$  Bright Data product. (We use Scraping Browser via WSS, plus the demand-index moat that wraps it.)

## Field: Project name

---

LongTail Scout

## Field: Tagline / short description (under 200 chars)

---

Apollo for the long tail. An AI agent that finds the small, local, nic



(190 chars — under the limit.)

## Field: Track / category

---

GTM Intelligence (Track 1)

## Field: Live demo URL

---

<https://longtailscout.com>

## Field: Source code URL

---

`https://github.com/dholzric/longtail-scout`

## Field: Video URL

---

[paste your YouTube unlisted URL here after upload]

## Field: Long description / "What does your project do?"

---

Paste this whole block verbatim:

LongTail Scout is a GTM-intelligence agent that surfaces the operators

### THE PROBLEM

GTM teams at vertical-SaaS companies (AccuLynx, ServiceTitan, Brightwh

Apollo and ZoomInfo can't help here. Their data was built around the L

### THE PRODUCT

Type a niche × city query: "roofing contractors in Houston". Get a ran

- Company name + homepage URL
- Hiring signals (open roles scraped from their careers page)
- Recent activity headlines pulled live from the web
- Size estimate
- A one-sentence sales angle generated from the scraped facts
- Every claim citation-linked back to the exact Bright Data fetch th

Per-row drill-down shows the underlying evidence. An on-page map plots

### WHY BRIGHT DATA IS THE SPINE

The agent runs a 3-phase pipeline on a single Cloudflare Worker:

1. DISCOVERY – LLM fires 4 parallel SERP queries via Brave Search AP
2. ENRICHMENT – for each top candidate, plain-fetch the homepage, pa

3. SYNTHESIS – DeepSeek ranks the operators and writes citation-group

#### THE KILLER DIFFERENTIATOR – NICHE RECON

The feature Apollo physically cannot replicate: paste your product des

Example output for "home services CRM for trades":

- electrical 100% Apollo-thin (every operator uses ServiceTitan b
- hvac 86% thin
- pool service 76% thin
- plumbing 50% thin

#### THE STACK

Cloudflare Worker (TypeScript, Wrangler) + Cloudflare KV (caching) + C

#### WHAT'S NEW VS. THE BASELINE

- Live scout in ~40 seconds end-to-end (was 8+ minutes before speed
  - Editorial / field-manual UI inspired by analog scientific journals
  - PER-OPERATOR INTELLIGENCE SUITE (v1.2-1.7), all Bright-Data-backed
    - Apollo-blind verification – a live site:linkedin.com/company s
    - Contact discovery – walks the contact/about pages for a real e
    - Decision-maker finder – the owner/founder + their LinkedIn /in
    - Signal radar – live third-party news (funding / expansion / le
    - "Act first" trigger feed – re-ranks the run by buying-signal s
    - Account-brief export – a one-click Markdown dossier (evidence
- The drill-down is now a full loop: find → prove it's Apollo-blind
- MCP server at /api/mcp (12 tools – scout, find\_businesses, demand\_
  - Live cost meter (Bright Data renders + LLM tokens in USD, streamin
  - Watchlist with daily cron-refreshed demand counts + Resend email d
  - Demand-index hostname-match geocoding (when Nominatim misses, we i
  - 118 unit tests including SSRF guards (worker + bridge), HMAC unsubs
  - Two passes of external security review (Codex), 10/10 findings clo

#### LIVE DEMO

<https://longtailscout.com> – see the Demo Notes section below for the p

#### SOURCE

<https://github.com/dholzric/longtail-scout>

## Field: How does it use Bright Data?

---

Bright Data Scraping Browser is the deep tier of every web call the ag

DISCOVERY tier 3 fallback: when Brave SERP + DDG miss a long-tail nich

ENRICHMENT JS fallback: most operator homepages render with plain HTTP

OPERATOR SCREENSHOTS: every operator's drill-down shows a homepage scr

The Bright Data bridge exposes /render, /serp, /screenshot – HTTP endp

Configured zones:

- lts\_browser (browser\_api): single zone handles SERP + render + scr

## Field: Inspiration / why this matters

---

GTM tooling has been about the same 100,000 enterprise accounts for a

The opportunity isn't a better Apollo. It's a parallel data layer that

That's the bet LongTail Scout makes.

## Field: Challenges we ran into

---

1. SPEED. Live scouts initially took 8+ minutes because Bright Data Sc serializes per-context renders, and our agent fired 25+ enrichment it. The fix was a three-layer cascading SERP tier (Brave → DDG HTML plain-fetch-only enrichment as the default (BD as opt-in quality ti parallel-with-stagger discovery instead of one-at-a-time tool calls scout time dropped from ~480s to ~40s.
2. JOB-TITLE HALLUCINATION. The careers-page parser was matching <li> phone-country-code dropdowns. The trace literally showed "Cook Isla parsed job title. Fixed with a blacklist regex (phone codes, emails "page-level hiring marker required before parsing list items" rule.
3. SECURITY. Two passes of external code review (Codex) found 10 issue the demand-index API was anonymous-accessible (moat exposure), the leaked every subscriber's email and webhook URL, the email unsubscr

dead path (CAN-SPAM violation), and – caught by the unit tests we w  
final security pass – an IPv6 SSRF bypass in the screenshot endpoint  
PRIVATE\_HOSTNAME\_RE only matched literal "fc00:" / "fd00:" / "fe80:  
addresses like [fd12:3456::1] slipped through). All 10 findings clo

4. THE NICHE RECON SIGNAL. The first version computed "Apollo-thinness  
businesses without a website" – but the demand index records SOME U  
every row (Booksy booking page, Boulevard widget, Yelp profile, Fac  
Google Maps redirect). So thinness was 0% for every niche, and the  
useless. The fix was redefining thinness as "% of businesses whose  
on a 50-platform blocklist" – booking platforms and social profiles  
as own-domain. Suddenly electrical contractors showed 100% Apollo-t  
operator using ServiceTitan booking) and the feature came alive

## Field: Accomplishments we're proud of

---

- 40-second live scouts. Real Bright Data calls, real LLM inference, c  
claim. The trace is visible to the user as it happens, with a live c  
ticking in cents.
- The Niche Recon feature. Reverses the GTM funnel – given a product d  
the verticals where Apollo is structurally blind. The killer demo mo  
physically cannot reproduce.
- The per-operator intelligence suite that closes the GTM loop: live A  
LinkedIn verification, contact discovery, decision-maker finder, sig  
one-click account brief – find -> prove -> reach the person -> why-n  
"prove it's Apollo-blind" and "find the actual person" steps are thi  
data vendor does.
- 118 unit tests covering the security-critical paths (worker + bridge  
unsubscribe tokens, PII redaction, Apollo-thinness detection) and ev  
The IPv6 SSRF bypass was caught by tests we wrote BEFORE shipping th  
exactly how unit tests should pay for themselves.
- A full live end-to-end burn-test of all 12 API endpoints AND all 12  
real Bright Data, which caught 3 real bugs (a same-zone-loopback 522  
a garbage-name parse, a cache-key leak) – all fixed and re-verified
- MCP server with 12 tools that any MCP-aware client (Claude Desktop,  
can drive directly. The same agent pipeline works as a website AND a  
for a separate LLM.

- An editorial / field-manual UI inspired by analog scientific journal serif, JetBrains Mono, cream-paper palette, paper-grain texture). The language reinforces the product's claim – this is reference material chatbot.
- The whole thing runs on Cloudflare Workers' free-tier-friendly stack custom domain, single KV namespace, one cron trigger. ~\$0.03 per sco external API spend

## Field: What we learned

---

- LLMs are dramatically better at structured reasoning when you GIVE t from a private index instead of asking them to estimate. "Roofing ha businesses" beats "roofing is a large vertical."
- Plain HTTP fetches recover ~60% of US-SMB homepages. The remaining 4 Bright Data's Browser API. That hybrid (plain-fetch first, BD on dem the right cost/speed/coverage trade-off.
- DDG SERP fallback is real and free. A regex against DDG's HTML page parseable results without any auth. Saved us during the ~24h Brave S rate-limit window.
- Cloudflare Workers + KV + Cron + Tunnels makes a real production dep achievable in a weekend. The frontend, the backend, the cron trigger custom domain, and the API tunnel for the private demand server – on account, no infra ops.



## Field: What's next

---

- CRM connectors (HubSpot + Salesforce one-click import, Apollo CSV reverse-import).
- "Always-on" mode: re-run the same niche x city weekly, surface new e to the demand index, fire Slack/email digests on delta > 0. (Already as cron + Resend; needs a UI to enable per-watch).
- Vertical-tuned prompts: roofing, HVAC, childcare, EMS-specific hiring signals so the LLM weights different evidence per market.
- Quality mode: opt-in BD-bridge enrichment that recovers the ~35% of

- candidates plain-fetch drops, in exchange for 3-5x scout latency.
- Browser MCP integration via Bright Data's own MCP server – once it's generally available, deprecate our playwright bridge and route direct

## Field: Tech tags

---

Bright Data, DeepSeek, Cloudflare Workers, TypeScript, Preact, Playwri



## Field: Team

---

Dan Holzrichter (solo)

## Field: Demo Notes (judge-facing, PRIVATE — not the public description)

---

These go in the submission's private notes / judge comments field, NOT the public description. They contain the demo password.

DEMO ACCESS

Open <https://longtailscout.com/?key=Piglet> – the ?key= captures the de into localStorage automatically and strips it from the URL on first lo

If the demo password prompt appears, paste: Piglet

WHAT TO TRY (in order of demo impact)

1. The default query is "roofing contractors in Houston". Click "RUN S watch the live trace + cost meter + inline map populate. ~40 second from a warm cache.
2. After results render, click any row to drill down. Every claim is f with a citation link back to the Bright Data fetch that produced it

3. Click the "🔍 don't know your niche? get recon →" chip above the q  
Paste a product description (try "home services CRM for trades – qu  
dispatch, mobile job sheets, payments"). Click "find niches →". You  
the top 5 long-tail verticals ranked by Apollo-thinness – this is t  
feature we built around our private 7M-business demand index. ~30-5  
cold cache, ~3s on a warm one.
4. The MCP server lives at <https://longtailscout.com/api/mcp> – try:  
`curl https://longtailscout.com/api/mcp` (lists 12 tools)  
See /mcp page for Claude Desktop config.
5. Map view (toggle at the top of the result table) shows operators as  
pins against a heat underlay of every other business in the demand  
that niche.

#### KNOWN LIMITS

- Quality vs. speed trade-off: live scouts use plain HTTP enrichment b  
(the "live" pace would otherwise be 3-4 minutes per scout via Bright  
Browser API serialization). Plain HTTP recovers ~60% of US-SMB homep  
the rest are dropped from the result set. The Bright Data Browser AP  
still used for SERP fallbacks, screenshots, and JS-heavy career page
- ?sample=1 mode returns cached results deterministically in ~140ms wi  
BD or LLM cost – useful if you want to demo the UX without burning c  
Example: <https://longtailscout.com/?q=roofing+contractors+in+Houston>

If anything breaks during your evaluation, please email [dholzric@gmail](mailto:dholzric@gmail.com)  
watching for judge traffic and can fix issues in minutes

